

SPECTRALMOON STUDIO · METHOD NOTE

The *Vocabulary* of a Directed Image

How we talk to image and video models so they stop guessing. Two systems, kept deliberately separate, and the order of operations that keeps the bill survivable.

FOR MARK, AI BUILDERS ×10

SUSANA BARRERO · SPECTRALMOON STUDIO

IMAGINATION IS THE TECHNOLOGY

One principle, and *everything* follows

A model does not feel. It executes statistical associations. So the entire craft reduces to a single move: **translate every feeling into named physics.**

The AI never makes a mistake. An under specified brief does.

"Make it cinematic" is not an instruction, it is a wish. Lens, f-stop, angle, light direction, Kelvin, material, palette: those are instructions.

VAGUE	DIRECTED
Good lighting	Key from a large softbox 45° camera right, silver fill camera left at quarter power, rim behind for edge separation
Beautiful sunset	Golden hour 6:30 PM, 200mm compression, backlit rim, warm 2800K, anamorphic flare
Luxury watch	Brushed steel case with micro scratches, macro, hard side light revealing texture
Make it cinematic	Deakins in Blade Runner 2049: amber and teal, atmospheric haze, strong directional light, geometric shadow

Specificity creates consistency. Vagueness creates chaos.

Two systems, kept *separate*

We run two complete methods. They are not rivals and they are not interchangeable. One is the container, the other is what goes inside it. The rule that matters: **never blend their vocabulary or structure inside a single prompt.** Pick the gate, stay in it.

	GENHQ	OHNEIS, THE VISUAL DIRECTOR METHOD
What it is	A production system	An art director's method
Question it answers	Did I specify everything?	Is this any good, and what exactly do I say to fix it?
What it gives you	Eight slots, the Motion Rule, character sheets, a fixed output format, the campaign steps	The technical language library, the layered prose order, visual diagnosis, the repair formula
Shape	A checklist	A vocabulary and an eye
Reach for it when	You need completeness and repeatability	You need execution detail, judgement, or recovery

THE ROUTING RULE

This is the piece most people never learn, and it is invisible when you get it wrong. The same knowledge is written two different ways depending on which model reads it.

WRITE IT AS	SEND IT TO	BECAUSE
A terse slot table, comma separated	Midjourney and still models	Terseness reads as control
Natural technical prose, full sentences	Kling, Seedance, Veo, Runway, Nano Banana Pro	Video models parse sentences. Comma soup reads to them as noise

GenHQ: the *completeness* system

Eight slots. Every image prompt, every video prompt, no exceptions and no silent drops. If a slot is left empty the model fills it for you, and it will not choose what you wanted. That is the whole argument for the checklist.

SLOT	WHAT IT DECIDES	IF YOU LEAVE IT BLANK
1 • Medium	Photograph, film still, illustration, 3D render	The model picks a physics engine at random
2 • Shot	Macro, close, medium, wide, establishing	You get a mid shot of everything
3 • Angle	Eye level, low, high, worm, bird, Dutch	Default eye level, all power drained
4 • Movement	Camera and subject motion. For stills write "static, locked off"	Video goes limp or drifts arbitrarily
5 • Focus	Shallow, deep, soft, rack. Always name what is sharp	The model chooses the hero, and it chooses wrong
6 • Subject	The specific who and what, doing what	Generic stock human
7 • Lighting	Source, direction, quality, Kelvin	Flat ambient, the look of nothing
8 • Colour palette	The named grade, and the strip image that locks it	Every frame drifts to its own world

The Motion Rule

Never leave both the subject and the camera static in a video prompt. One of them must move or the scene dies on screen.

THE OUTPUT FORMAT IS PART OF THE SYSTEM

Every prompt is delivered in the same four parts, in this order, so it can be audited and pasted without editing: the **eight slot table**, the **prose prompt**, the **negatives**, and a **specs block** naming tool, model, duration, aspect ratio, resolution and reference images.

Negatives are written as "no X", never "avoid X". Bare tokens enforce weakly. This holds inside tool native fields too, including Midjourney's own negative parameter.

Ohneis: the *execution* order

Where GenHQ guarantees nothing is missing, Ohneis governs how it is actually said. This is the prose dialect: flowing sentences that still name every specification. Never keyword soup, never vague description. Written in a fixed order, because the order is what prevents hallucination.

- 01 **Camera.** Sensor, lens in mm with its character, aperture with the reason for it, ISO, the overall look.
- 02 **POV and coordinates.** Camera height and what it faces, then every element placed by thirds: left, centre or right third, lower, middle or upper band. State what sits behind camera, frame left, frame right, foreground. Name the leading lines and the depth layers from foreground through midground to background.
- 03 **Lighting.** Source, direction, quality, time of day, and the Kelvin value spoken aloud.
- 04 **Mood and palette.** A named palette, plus the clichés you are explicitly forbidding.
- 05 **Subject.** The specific hero. This is the one swappable line in the whole prompt.
- 06 **Material reality.** True surfaces and true textures.
- 07 **Style and post.** An art director or photographic reference, the grade, the grain.

Coordinates are the anti hallucination layer

If a door, a wall or a building appears where it should not, the coordinate layer was incomplete. The model holds no map of your scene. It knows only the frame. So never write route language like "walking in from the pool", which is meaningless to it. Write pure frame geometry: what sits in which third, at what angle, and what is off frame right.

Negatives are baked in as positives, woven where they belong, never dumped in a block at the end. Write "carefully avoiding golden or orange grading", not a list of banned words.

The camera *vocabulary*

Four tables carry most of the control. Pick one value per dimension, deliberately, and know why.

FOCAL LENGTH, WHICH IS PERSPECTIVE

RANGE	CHARACTER	REACH FOR IT WHEN
8 to 16mm	Extreme distortion, drama	Interiors, edgy and experimental
16 to 35mm	Environmental context	Street, real estate, scene setting
35 to 50mm	Natural human vision	Documentary, versatile
50 to 85mm	Compression, isolation, flattering	Portraits
85 to 135mm	Beautiful bokeh	Fashion, professional portraits
135 to 300mm	Extreme compression	Tight candids, wildlife

APERTURE, WHICH IS DEPTH

F-STOP	EFFECT
f/1.2 to f/1.8	Dreamlike bokeh, low light, mood
f/2.0 to f/2.8	Portrait sweet spot, sharp subject
f/4 to f/5.6	Balanced, groups, travel
f/8 to f/11	Landscape and architectural sharpness
f/16 to f/22	Everything in focus, sunstars

ANGLE, WHICH IS POWER

ANGLE	FEELING
Worm's eye	Maximum drama and power
Low	Dominance, heroic, architectural strength
Eye level	Neutral, documentary honesty

High	Vulnerability, overview
Bird's eye	Pattern, abstraction
Dutch	Tension, unease, energy

LIGHT, WHICH IS FOUR DECISIONS AND NOT ONE

Source: window, studio strobe, neon, candle, on camera flash, street ambient. **Direction:** front, side, back, rim. **Kelvin:** warm tungsten 2700 to 3200K, neutral daylight 5600K, cool LED 6500K. **Quality:** hard for sharp shadow and texture, soft for flattering evenness, directional for sculptural shape, diffused for clean commercial, backlit for silhouette and separation.

Colour is an *image*, not a sentence

This is the highest leverage habit on the list, and the one almost nobody has.

A model cannot hold colour consistent across scenes on its own. Generate six frames independently and you get six slightly different worlds. The set stops reading as one film and starts reading as a collection.

Hex written into a prompt is a suggestion. A palette strip passed as an image is a lock.

So we lock the palette as hex, then **render those hexes into a horizontal strip PNG and pass that strip into every single frame** as an explicit colour reference. One shared input, and the whole set grades to one world.



The strip this document is graded to, shown the way it would be passed to a model: an image, not a description. One hue at 75°, one saturation, six steps of lightness. Note there is no pure black. The darkest value is a deep sage, which keeps prints and grades from going dead.

THE FIVE INPUT STACK, ONE JOB EACH

INPUT	LOCKS
1 · Location render	Lighting, atmosphere, overall colour feel
2 · Composition sketch	Framing, and where everyone stands
3 · Character reference	Face, hair, wardrobe
4 · Second character	The same, held separate so the two do not blend
5 · Palette strip	The exact grade, shared across every scene

Separating the inputs means each one carries a single job, so nothing fights. Most drift comes from asking one image to do three things at once.

Characters, and how they *drift*

Character consistency is its own discipline, and it has one artifact at the centre of it.

The character sheet is a single composite turnaround: five angles in one image, generated in one fire. Never five separate generations. Five separate fires give you five subtly different people, which is the exact problem the sheet exists to solve.

That sheet then feeds every subsequent generation involving that character, **at any scale, even when the figure is tiny or far away.** A distant figure described in text will drift: wrong dress, wrong hair, wrong build, and you pay to discover it.

For a character who recurs across a whole body of work, promote the sheet into a **Soul ID**, a trained character lock inside the video tool, so identity holds without re-passing references every time.

Three traps we paid to learn

Never pass the full turnaround when placing someone into a scene. The model averages the poses and keeps the dominant facing direction, ignoring your text. Crop the single view whose pose and direction match the shot, flip it if needed, and pass that alone.

Never reposition a character using a scene reference that already contains them.

The model anchors to the baked in pose, or adds a duplicate. Generate a clean empty plate first, then place the cropped character into it.

Never re-describe a character who is locked by reference. Say "the figure" and let the image carry identity. Repeating appearance in text fights the reference, and stacking body words raises the odds of a content filter trip. Spend the prompt on motion, action and environment instead.

When a frame *rots*

Feed a generated image back in as a reference enough times and damage compounds: smear, plastic skin, melted detail. Re-prompting the damaged frame only bakes the damage deeper.

Think of an image as two layers. **Low frequency** is information that does not repeat: geometry, layout, the bones. **High frequency** is the surfaces that do repeat: skin, fabric, grain, the texture. When quality degrades it is almost always the high frequency layer rotting. The bones are usually still fine.

So do not re-roll. Strip the image to an outline sketch, which throws away the ruined texture and keeps the structure, then render brand new surfaces onto those clean bones. You are not repairing pixels, you are rebuilding skin on a good skeleton.

The human stays *in the loop*

This pipeline is agentic. There is a goal, tools get called, state is carried between steps, and the next move is chosen from that state. But every generation passes a human gate before it fires, and that is a deliberate design decision rather than an unfinished one.

| *Autonomy is not free when every action costs credits.*

Generation is the one class of tool call where a wrong decision is billed immediately and cannot be undone. So the machine plans, routes, holds state and prepares the exact prompt, then stops and shows its work. A person says go. Silence is never approval.

What the machine owns: reading the storyboard, choosing the tool, assembling the prompt to spec, tracking which frame feeds which, naming and filing every output, noticing when something drifted.

What the person owns: the go ahead, the aesthetic verdict, and the money.

The order of *operations*

Compressed to its spine, this is how a piece gets made. The sequence is the cost control.

1. **Storyboard the whole idea first.** The entire story, never one frame at a time. Frame by frame invention is where continuity dies.
2. **Lock the palette** as hex, then render it as a strip image so it can be passed, not described.
3. **Lock the characters.** One composite sheet, five angles, one fire. Promote to a Soul ID if the character will recur.
4. **Build the stills in the cheap tool.** Stills are where you iterate, because iteration there costs almost nothing.
5. **Validate cheap, then commit expensive.** Every first pass runs in the fast or draft mode. Standard and high quality modes are only ever spent on something already approved. This applies to stills and to video alike.
6. **Approve the still** before any video is considered. An unapproved still becomes an expensive unapproved video.
7. **Chain the video,** each shot starting from the previous shot's true last frame, so the walk is continuous rather than a set of cuts.
8. **Enhance last.** Upscaling and frame rate work happen once, at the end, on the approved take only.
9. **Never generate video in the expensive tool.** We learned this at a cost of one five second clip billed at roughly forty seven times the price of a still.

Reference first, always. When a reference image is available, feed it. Image to image beats text to image on nearly every axis: style, lighting, identity, composition.

Compiled from the GenHQ campaign system, the Ohneis Visual Director method, and the colour consistency work of @zhuchkin.

The mistakes documented here are our own.

Susana Barrero · spectralmoonstudio.com · 2026